

Bio-Informatics

Search Space Reduction in Genome-Wide Association Studies (GWAS) Using Machine Learning

Proposal

Team Members:

258247B – Keirishan B

258284J - Rathnapala G.D.S.K

Table of Contents

1. Background and Introduction	3
1.1 Genome-Wide Association Studies (GWAS)	3
1.2 The Search Space Problem in GWAS	3
1.3 Role of Machine Learning in Search Space Reduction	4
2. Literature Review (Key Papers)	5
Paper 1: LASSO for GWAS Feature Selection	5
Paper 2: Random Forest for SNP Prioritization	5
Paper 3: Relief-Based Feature Selection	5
Paper 4: Deep Learning for Representation Learning	6
Paper 5: Biological Knowledge-Guided Filtering	6
3. Problem Statement	8
4. Proposed Method	9
Stage 1: LD-Based Preprocessing	9
Stage 2: Weak-Effect Preserving Screening	9
Stage 3: Interaction-Aware Feature Selection	9
Stage 4: Representation Learning (Optional Extension)	9
Stage 5: Interpretability and Biological Validation	9
5. Available Datasets for Search Space Reduction in GWAS	10
5.1 UK Biobank	10
5.2 1000 Genomes Project	10
5.3 dbGaP (Database of Genotypes and Phenotypes)	10
5.4 WTCCC (Wellcome Trust Case Control Consortium)	10
5.5 GWAS Catalog	11
Conclusion	11

1. Background and Introduction

1.1 Genome-Wide Association Studies (GWAS)

Genome-Wide Association Studies (GWAS) are widely used to identify genetic variants associated with complex diseases and traits. In a typical GWAS, researchers analyze hundreds of thousands to millions of Single Nucleotide Polymorphisms (SNPs) across the genome in large cohorts to detect statistical associations between genetic variants and phenotypes.

GWAS has successfully identified thousands of loci associated with diseases such as diabetes, cardiovascular disorders, cancer, and psychiatric conditions. However, the method faces several computational and statistical challenges:

- High dimensionality ($p \gg n$ problem)
- Multiple testing burden
- Weak effect sizes of individual variants
- Complex interactions among loci (epistasis)
- Large computational cost

The human genome contains millions of SNPs, but only a small fraction are truly associated with any given phenotype. Therefore, efficiently reducing the search space before association testing is crucial.

1.2 The Search Space Problem in GWAS

In GWAS, the “search space” refers to the set of genetic variants (SNPs) and their potential interactions considered for analysis.

For example:

- 1 million SNPs
- Pairwise interaction testing approximately 10^{12} combinations
- Higher order interactions computationally infeasible

Brute-force testing:

- Increases false positives
- Requires very strict multiple testing corrections (e.g., Bonferroni correction)
- Reduces statistical power
- Is computationally expensive

Therefore, search space reduction aims to:

- Filter irrelevant SNPs before analysis
- Select informative features
- Prioritize biologically meaningful regions
- Reduce dimensionality
- Improve statistical power
- Enable interaction detection

Machine learning has emerged as a powerful tool to address this challenge.

1.3 Role of Machine Learning in Search Space Reduction

Machine learning helps reduce the GWAS search space through:

1. Feature selection methods
2. Regularization techniques
3. Embedded selection within predictive models
4. Biological knowledge-guided filtering
5. Representation learning

Unlike traditional univariate filtering (e.g., p-value thresholding), machine learning methods can:

- Capture nonlinear effects
- Model multivariate interactions
- Handle correlated SNPs
- Improve generalization

Thus, ML-based search space reduction improves both computational feasibility and biological discovery.

2. Literature Review (Key Papers)

Paper 1: LASSO for GWAS Feature Selection

Contribution:

This study applied LASSO (Least Absolute Shrinkage and Selection Operator) to GWAS data.

Key Idea:

LASSO introduces L1 regularization, shrinking irrelevant SNP coefficients to zero, effectively performing feature selection.

Strengths:

- Handles high-dimensional data
- Performs automatic variable selection
- Reduces overfitting

Limitations:

- Struggles with highly correlated SNPs (linkage disequilibrium)
- May select only one SNP from a correlated block

Paper 2: Random Forest for SNP Prioritization

Contribution:

Applied Random Forest (RF) to identify important SNPs.

Key Idea:

Random Forest ranks SNPs using feature importance measures.

Strengths:

- Captures nonlinear relationships
- Detects interactions implicitly
- Robust to noise

Limitations:

- Biased toward SNPs with more categories
- Difficult to interpret biologically
- Computationally expensive for millions of SNPs

Paper 3: Relief-Based Feature Selection

Contribution:

Improved Relief-based algorithms for detecting interactions in genetic data.

Key Idea:

Relief methods evaluate SNP importance by comparing nearest neighbors.

Strengths:

- Effective in detecting epistatic interactions
- Works without explicit interaction modeling

- Reduces search space for higher-order analysis

Limitations:

- Sensitive to noise
- Parameter tuning required

Paper 4: Deep Learning for Representation Learning

Contribution:

Surveyed deep learning methods for genomic data.

Key Idea:

Neural networks can learn compressed representations of genomic features, effectively reducing dimensionality.

Strengths:

- Captures complex nonlinear patterns
- Learns hierarchical representations
- Integrates multiple data types

Limitations:

- Requires large sample sizes
- Limited interpretability
- Risk of overfitting in small cohorts

Paper 5: Biological Knowledge-Guided Filtering

Contribution:

Proposed grouping SNPs into genes or pathways before testing.

Key Idea:

Reduce search space using biological priors such as gene boundaries, pathway databases, and functional annotations.

Strengths:

- Reduces multiple testing burden
- Biologically interpretable
- Improves power

Limitations:

- Depends on prior knowledge
- May miss novel loci

Summary of Literature

Method	Reduction Strategy	Captures Interactions	Interpretability
LASSO	Embedded feature selection	Limited	Moderate
Random Forest	Feature importance ranking	Yes	Low
Relief	Interaction-aware filtering	Yes	Moderate
Deep Learning	Representation learning	Yes	Low
Pathway-based	Knowledge-guided grouping	Indirect	High

Existing approaches reduce dimensionality but face trade-offs between interpretability, scalability, and statistical power. Many methods either fail to scale to biobank-level data or struggle to detect weak interaction effects without excessive computational cost.

3. Problem Statement

Despite significant progress in applying machine learning to GWAS, an important open problem remains: how to effectively reduce the search space for large-scale GWAS while preserving weak marginal effects and complex epistatic interactions in a computationally scalable and interpretable manner.

Current methods exhibit several limitations:

1. Univariate filtering methods remove SNPs based solely on marginal p-values, potentially discarding variants involved only in interactions.
2. Regularized regression methods such as LASSO struggle with correlated SNPs due to linkage disequilibrium and may arbitrarily select one SNP from a correlated block.
3. Tree-based ensemble methods like Random Forest can capture nonlinearities but do not scale efficiently to millions of SNPs without aggressive pre-filtering.
4. Deep learning models require large sample sizes, lack interpretability, and risk overfitting in moderate-sized GWAS datasets.
5. Knowledge-guided filtering depends heavily on existing biological annotations and may fail to identify novel genetic mechanisms.

Therefore, there is a need for a hybrid, scalable, and interaction-aware search space reduction framework that:

- Preserves SNPs involved in weak or purely interaction-based effects
- Handles linkage disequilibrium structures
- Scales to biobank-level datasets
- Maintains interpretability for biological validation

This gap motivates the development of an improved machine learning–based search space reduction strategy tailored for modern large-scale GWAS datasets.

4. Proposed Method

To address the identified problem, we propose a multi-stage hybrid search space reduction framework combining statistical filtering, interaction-aware feature selection, and representation learning.

Stage 1: LD-Based Preprocessing

- Perform quality control (minor allele frequency filtering, missingness filtering, Hardy–Weinberg equilibrium testing).
- Group correlated SNPs into linkage disequilibrium (LD) blocks.
- Select representative SNPs from each LD block to reduce redundancy while preserving genetic information.

This step reduces dimensionality while accounting for correlation structure.

Stage 2: Weak-Effect Preserving Screening

Instead of strict p-value thresholding, apply relaxed statistical screening combined with effect-size ranking to retain SNPs with weak but potentially meaningful effects.

- Use marginal regression with false discovery rate (FDR) control.
- Retain top-k SNPs per chromosome or LD block.

This ensures interaction-relevant SNPs are not prematurely discarded.

Stage 3: Interaction-Aware Feature Selection

Apply an interaction-sensitive machine learning method such as:

- Relief-based algorithms, or
- Gradient Boosted Trees with interaction depth control.

This stage identifies SNPs that contribute to predictive performance through nonlinear or interaction effects.

Stage 4: Representation Learning (Optional Extension)

Use an autoencoder or sparse neural network to learn compressed genomic representations.

- Input: filtered SNP matrix
- Output: lower-dimensional embedding
- Apply clustering or feature importance scoring on learned representation

This stage further compresses dimensionality while preserving higher-order structure.

Stage 5: Interpretability and Biological Validation

To ensure biological relevance:

- Map selected SNPs to genes and pathways.
- Compare with known associations from GWAS Catalog.
- Perform pathway enrichment analysis.

This hybrid framework integrates statistical rigor, machine learning power, and biological interpretability.

5. Available Datasets for Search Space Reduction in GWAS

5.1 UK Biobank

- Approximately 500,000 individuals
- Genotype and phenotype data
- Wide range of diseases and traits
- Imputed SNP data (tens of millions of variants)

Advantages:

- Large sample size
- Multiple phenotypes
- Suitable for ML models

Limitation:

- Requires application for access

5.2 1000 Genomes Project

- Whole-genome sequencing data
- Diverse populations
- Approximately 2,500 individuals

Advantages:

- High-quality variant reference
- Useful for simulation studies
- Population structure analysis

Limitation:

- Small sample size for disease association studies

5.3 dbGaP (Database of Genotypes and Phenotypes)

- NIH repository
- Contains multiple GWAS datasets
- Disease-specific cohorts

Limitation:

- Controlled access required

5.4 WTCCC (Wellcome Trust Case Control Consortium)

- Classic GWAS dataset
- Multiple common diseases
- Approximately 14,000 cases and controls

Advantages:

- Frequently used in benchmarking ML methods
- Well-documented

5.5 GWAS Catalog

- Curated list of published GWAS results
- Summary statistics
- SNP–trait associations

Advantages:

- Useful for validation
- Can guide feature prioritization
- Enables meta-analysis

Limitation:

- No individual-level genotype data

Conclusion

Search space reduction remains a central challenge in GWAS due to extreme dimensionality, computational constraints, and complex genetic architectures. While existing machine learning methods provide useful tools, no single approach fully addresses scalability, interaction detection, and interpretability simultaneously.

The proposed hybrid framework aims to combine statistical screening, LD-aware reduction, interaction-sensitive learning, and biological validation into a coherent pipeline suitable for large-scale genomic studies. This approach has the potential to improve computational efficiency while preserving biologically meaningful genetic signals.